

# Eye Movements

Daniela Pamplona

U2IS - ENSTA - IPParis

ecampus moodle: MI210 - Modèles neuro-computationnels de  
la vision (P4 - 2020-21)

[daniela.pamplona@ensta.fr](mailto:daniela.pamplona@ensta.fr)


# Contents

1. Major types of eye movements
2. What triggers saccades?
3. Reinforcement learning
  1. Markov decision processes
  2. Q learning with Monte Carlo and SARSA
4. Eye movements to learn to solve visually guided tasks

# Contents

1. Major types of eye movements
2. What triggers saccades?
3. Reinforcement learning
  1. Markov decision processes
  2. Q learning with Monte Carlo and SARSA
4. Eye movements to learn to solve visually guided tasks

# 3 groups of eye movements



# 3 groups of eye movements

- Ductions: isolated movements of a single eye




Elevation



Adduction



Intorsion



# 3 groups of eye movements

- Ductions: isolated movements of a single eye
- Vergences: vergence movements are mirror image movements, being equal and opposite





Elevation



Adduction



Intorsion



# 3 groups of eye movements

- Ductions: isolated movements of a single eye
- Vergences: vergence movements are mirror image movements, being equal and opposite
- **Versions:** conjugate gaze movements of both eyes





Elevation



Adduction



Intorsion



# Major types of versions movements

Image stabilization during body movements  
(involuntary)



# Major types of versions movements

Image stabilization during body movements (involuntary)

a) Vestibulo-ocular reflex: compensate head movement to keep object in the fovea, max velo: 350 deg/s



b) Optokinetic reflex: compensate large motions on the visual field, max velo:??



# Major types of versions movements


Image tracking/search of objects,  
**acquiring information** (voluntaty)

# Major types of versions movements

Image tracking/search of objects,  
**acquiring information** (voluntaty)

a) Smooth pursuit: follow continuously  
a moving object, max velo:  
100deg/sec

**b) Saccades:** bring object to fovea,  
jerky and abrupt, max velo:  
900deg/sec



# Major types of eye movements

Image tracking/search of objects,  
**acquiring information** (voluntary)

a) Smooth pursuit: follow continuously  
a moving object, max velo:  
100deg/sec

b) **Saccades**: bring object to fovea,  
jerky and abrupt, max velo:  
900deg/sec

Why do we need to bring objects to the fovea?



# Why do we need to bring objects to the fovea?




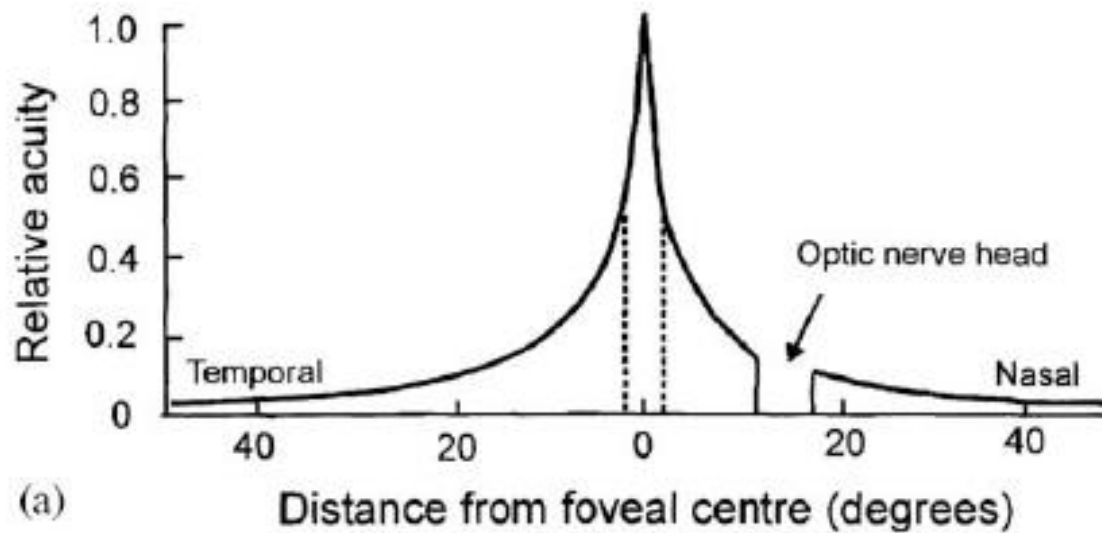


Fig. 20. Graph to show rod and cone densities along the horizontal meridian.



# Why do we need to bring objects to the fovea?



# Where is the oculomotor control center?




Figure 2.1 Brain areas that control eye movements

Eye movements and attention - Daniela Pamplona

# Where is the oculomotor control center?




Figure 2.1 Brain areas that control eye movements  
Eye movements and attention - Daniela Pamplona



# Contents


1. Major types of eye movements
- 2. What triggers saccades?**
3. Reinforcement learning
  1. Markov decision processes
  2. Q learning with Monte Carlo and SARSA
4. Eye movements to learn to solve visually guided tasks

# What triggers saccades?

# What triggers saccades?

## Models of attention

1. Bottom up theories
2. Top down theories
3. Mixture



# What triggers saccades?

## Models of attention


### 1. Bottom up theories

- 1) Saliency
- 2) Context

### 2. Top down theories

1. Memory
2. Task

### 3. Mixture



# Bottom up: Saliency

Stimulus




# Bottom up: Saliency

Stimulus



Gazing heat map




# Bottom up: Saliency


Stimulus



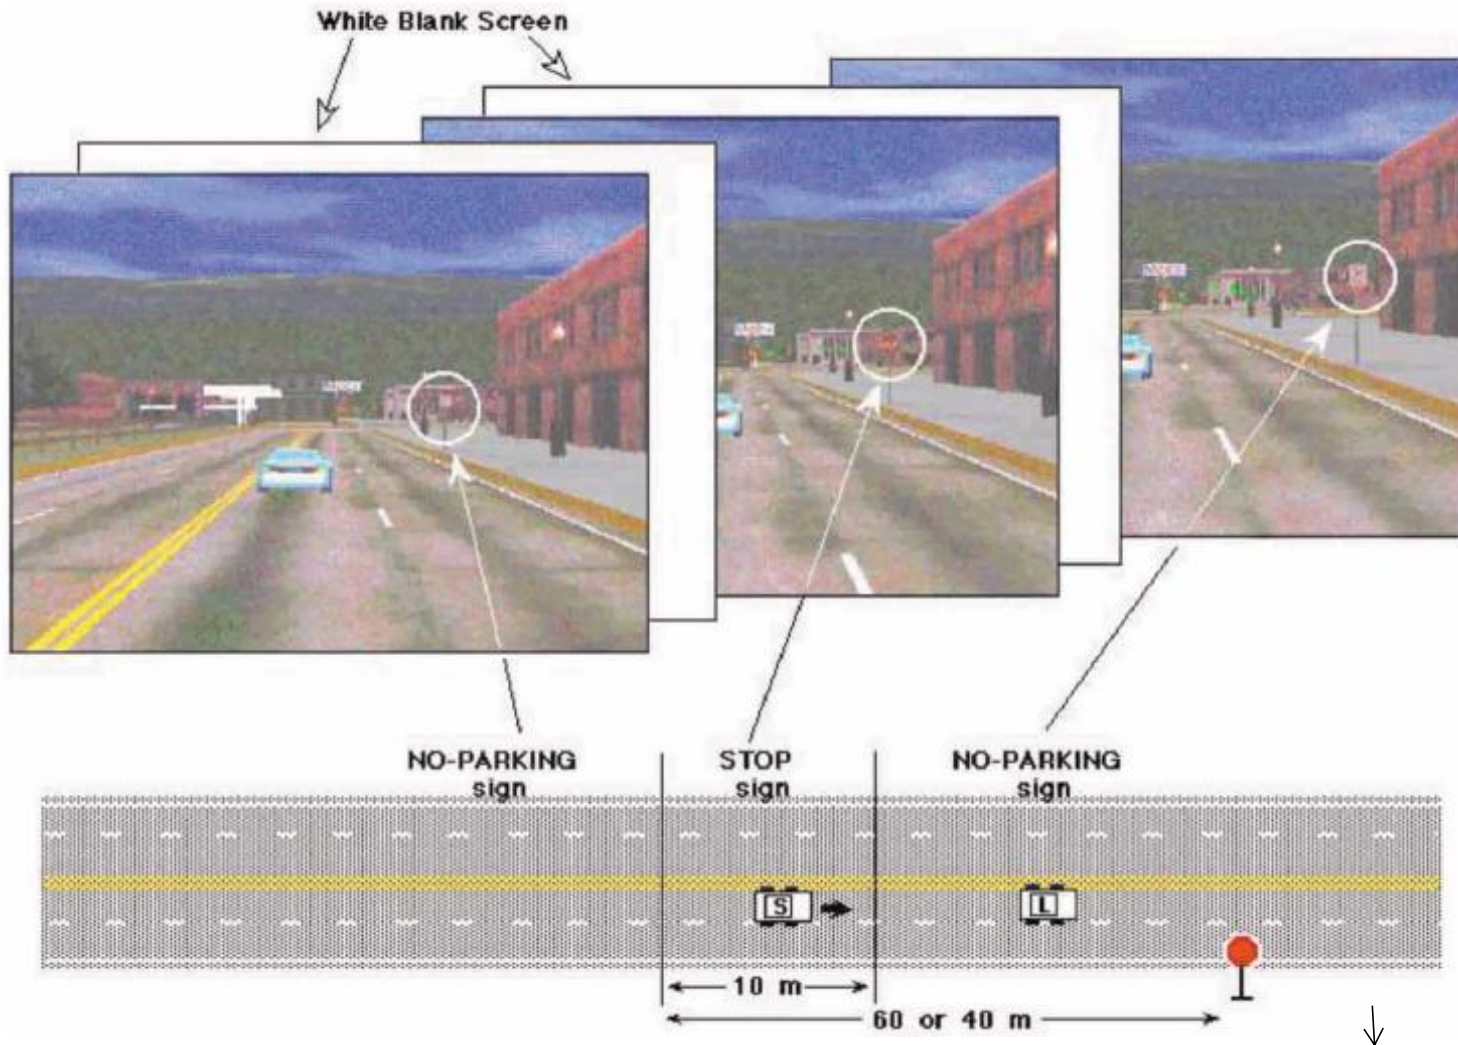
Gazing heat map



Saliency based on the responses to ICA


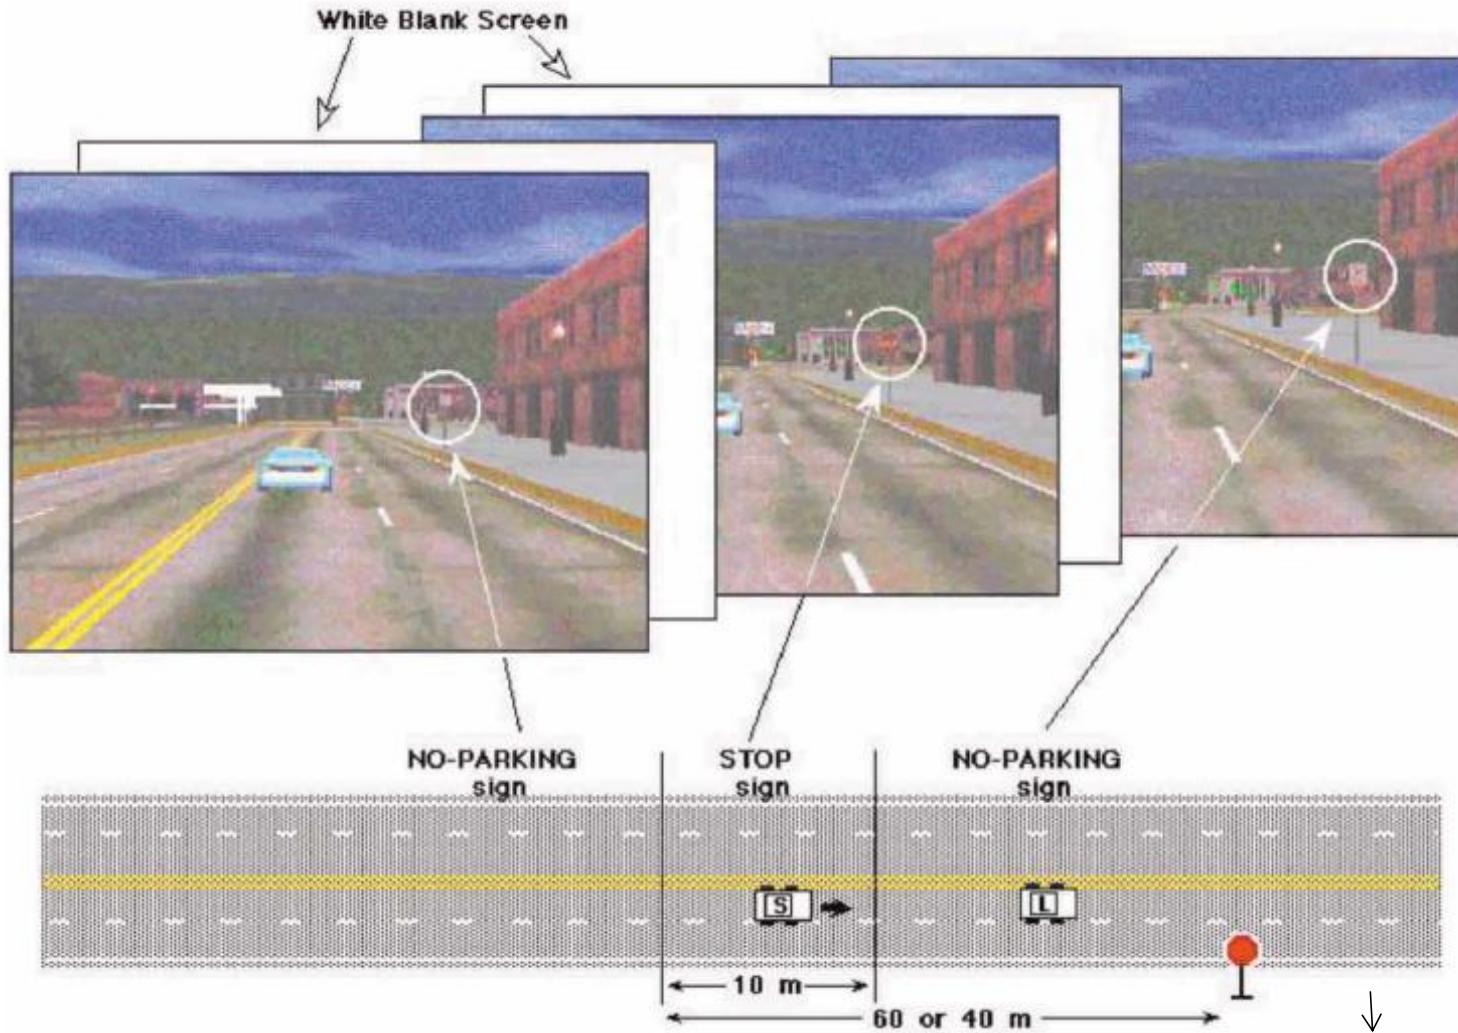


# Bottom up: Context





# Bottom up: Context



# Top down: Memory

Target



# Top down: Memory

Target



1<sup>st</sup> search



# Top down: Memory

Target



1<sup>st</sup> search



2<sup>nd</sup> search



# Top down: Memory

Target




1<sup>st</sup> search




2<sup>nd</sup> search



3<sup>rd</sup> search



# Top down: Task




Tasks/behaviors:

1. avoid obstacle
2. sidewalk following
3. pick litter




# Top down: Task




Tasks/behaviors:

1. avoid obstacle
2. sidewalk following
3. pick litter




# Top down: Task

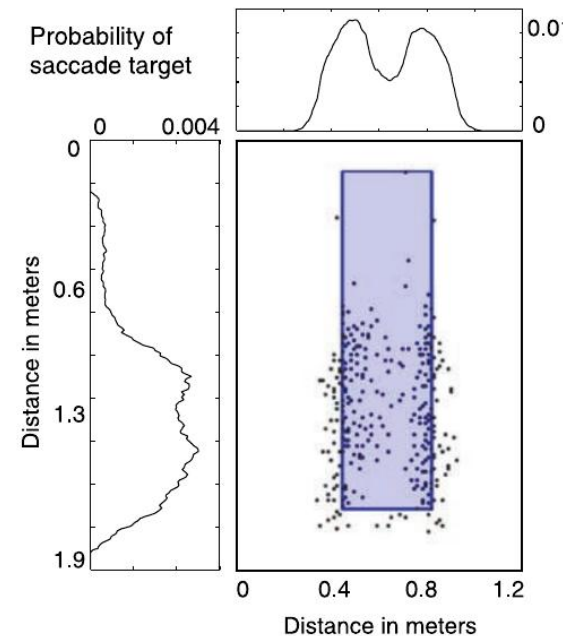


Tasks/behaviors:

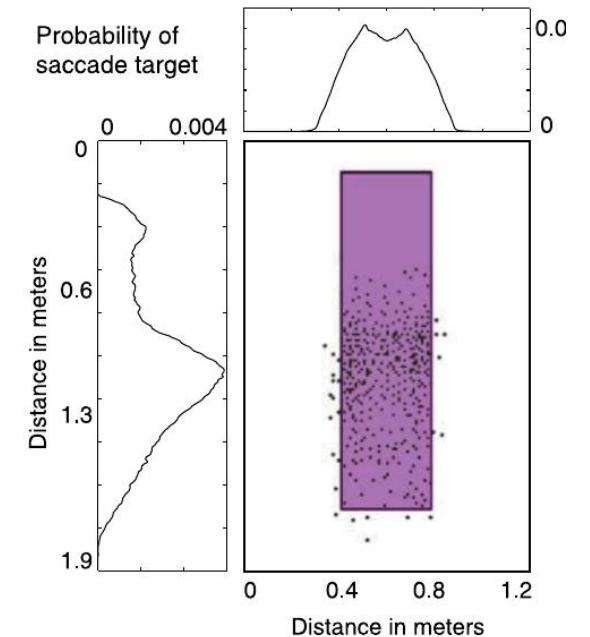
1. avoid obstacle
2. sidewalk following
3. pick litter



Gaze distribution on obstacles in the “avoid” condition




Gaze distribution on litter in the “pickup” condition







# Mixed



- Task: to make a PBJ sandwich
- 50% of objects were irrelevant (scotch tape, forks, etc)
  - before task started: 52% of time looking at irrelevant objects
  - when task started: 18% of time looking at irrelevant object

# Mixed

- Task: to make a PBJ sandwich
- 50% of objects were irrelevant (scotch tape, forks, etc)
  - before task started: 52% of time looking at irrelevant objects
  - when task started: 18% of time looking at irrelevant object




# Break!


# Contents

1. Major types of eye movements
2. What triggers saccades?
3. Reinforcement learning
  1. Are eye movements learned?
  2. Markov decision processes
  3. Q learning with Monte Carlo and SARSA
4. Eye movements to learn to solve visually guided tasks

# Learning eye movements for playing cricket




# Learning eye movements for playing cricket




Gaze and ball  
directions (angles  
from head)


Weak amateur



Time (s)




Time (s)




Time (s)

# Learning eye movements for playing cricket




Gaze and ball  
directions (angles  
from head)

Weak amateur




Time (s)

Good amateur




Time (s)




Time (s)

# Learning eye movements for playing cricket




Gaze and ball  
directions (angles  
from head)

Weak amateur




Time (s)

Good amateur



Time (s)


Professional



Time (s)



# From the shortest path problem to reinforcement learning




Goal to find the shortest path from A to F without repetitions

- $S \equiv$  set of states: A, B, C...
- $A \equiv$  set of actions: turn  $90^\circ$ ,  $45^\circ$ ,  $-30^\circ$
- $V \equiv$  distance between states
- $R \equiv$  reward the sum of all neg distances.

The optimal action sequence can be calculated by Dijkstra's algorithm

# From the shortest path problem to reinforcement learning




Goal to find the shortest path from A to F without repetitions

- $S \equiv$  set of states: A, B, C...
- $A \equiv$  set of actions: turn  $90^\circ$ ,  $45^\circ$ ,  $-30^\circ$
- $V \equiv$  ????
- $R \equiv$  reward the sum of all neg distances given

How do we solve this problem?

# From the shortest path problem to reinforcement learning




Goal to find the shortest path from A to F without repetitions

- $S \equiv$  set of states: A, B, C...
- $A \equiv$  set of actions: turn  $90^\circ$ ,  $45^\circ$ ,  $-30^\circ$
- $V \equiv$  ????
- $R \equiv$  reward the sum of all neg distances


How do we solve this problem?

# From the shortest path problem to reinforcement learning: Brute force

Episode 1: A->C->F:  $R(\text{Ep.1}) = -13$




# From the shortest path problem to reinforcement learning: Brute force



Episode 1: A->C->F:  $R(\text{Ep.1}) = -13$

Episode 2: A->B->C->F:  $R(\text{Ep.2}) = -14$

# From the shortest path problem to reinforcement learning: Brute force




Episode 1: A->C->F:  $R(\text{Ep.1}) = -13$

Episode 2: A->B->C->F:  $R(\text{Ep.2}) = -14$   
 $V(A,B)+V(B,C)=V(A,C) -1$

Episode 3: A->B->D->C->F:  $R(\text{Ep.3}) = -16$

...


# From the shortest path problem to reinforcement learning: Brute force



Problems with brute force solution:

- The number of trajectories grows exponentially
- The number of states might be infinite
- The reward might be stochastic

# From the shortest path problem to reinforcement learning




One step reward matrix

-	A	B	C	D	E	F
A	-	-3	-5	-9	-	-
B	-3	-3	-4	-	-7	-
C	-5	-4	-	-2	-6	-8
D	-9	-	-2	-	-2	-2
E	-	-7	-6	-2	-	-5
F	-	-	-8	-2	-5	-



# From the shortest path problem to reinforcement learning



One step reward matrix

-	A	B	C	D	E	F
A	-	-	-5	-9	-	-
B	-	-	-3	-	-7	-
C	-5	-4	-	-2	-6	-8
D	-9	-	-2	-	-2	-2
E	-	-7	-6	-2	-	-5
F	-	-	-8	-2	-5	-

IMMEDIATE REWARD,  
NOTHING ABOUT THE TOTAL REWARD

# RL with MDP: main functions

- **Expected discounted reward:** ballance between earlier and later rewards.

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \quad 0 \leq \gamma \leq 1,$$

# RL with MDP: main functions

- **Expected discounted cumulative reward:**

ballance between earlier and later rewards.



$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \quad 0 \leq \gamma \leq 1,$$

- **value of a state-action under a policy  $\Pi$ :** expected discounted cumulative reward starting from that state and taking that action

$$q_{\pi}(s, a) \doteq \mathbb{E}_{\pi}[G_t \mid S_t = s, A_t = a] = \mathbb{E}_{\pi}\left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a\right]$$

# From the shortest path problem to reinforcement learning

- Q-matrix



-	A	B	C	D	E	F
A	-	20			-	-
B				-		-
C			-			
D		-		-		
E	-				-	
F	-	-				-

# RL with MDP: main functions

- **Expected discounted return:** ballance between earlier and later rewards.

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \quad 0 \leq \gamma \leq 1,$$

- **value of a state-action under a policy  $\Pi$ :** expected discounted return starting from that state and taking that action

$$q_{\pi}(s, a) \doteq \mathbb{E}_{\pi}[G_t \mid S_t = s, A_t = a] = \mathbb{E}_{\pi}\left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a\right]$$

- **Greedy policy:**

$$\Pi^*(s) = \operatorname{argmax}_a Q(s, a)$$

# RL with MDP

- <https://www.youtube.com/watch?v=bHeeaXgqVig>

# on-policy: Monte Carlo method (extension to brute force)

Monte Carlo ES (Exploring Starts), for estimating  $\pi \approx \pi_*$

Initialize, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$ :

$Q(s, a) \leftarrow \text{arbitrary}$

$\pi(s) \leftarrow \text{arbitrary}$

$Returns(s, a) \leftarrow \text{empty list}$

Repeat forever:

Choose  $S_0 \in \mathcal{S}$  and  $A_0 \in \mathcal{A}(S_0)$  s.t. all pairs have probability  $> 0$

Generate an episode starting from  $S_0, A_0$ , following  $\pi$

For each pair  $s, a$  appearing in the episode:


$G \leftarrow$  the return that follows the first occurrence of  $s, a$

Append  $G$  to  $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

For each  $s$  in the episode:

$\pi(s) \leftarrow \arg \max_a Q(s, a)$



$$\pi_0 \xrightarrow{\text{E}} q_{\pi_0} \xrightarrow{\text{I}} \pi_1 \xrightarrow{\text{E}} q_{\pi_1} \xrightarrow{\text{I}} \pi_2 \xrightarrow{\text{E}} \dots \xrightarrow{\text{I}} \pi_* \xrightarrow{\text{E}} q_*,$$

# on-policy: Monte Carlo method (extension to brute force)

## Monte Carlo ES (Exploring Starts), for estimating

Initialize, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$ :

$Q(s, a) \leftarrow \text{arbitrary}$

$\pi(s) \leftarrow \text{arbitrary}$

$\text{Returns}(s, a) \leftarrow \text{empty list}$

Repeat forever:

Choose  $S_0 \in \mathcal{S}$  and  $A_0 \in \mathcal{A}(S_0)$

Generate an episode

For each pair  $s, a$ :

$G \leftarrow \text{the return from state } s \text{ after action } a$



Append  $G$  to  $\text{Returns}(s, a)$

$Q(s, a) \leftarrow \text{average}$

For each  $s$  in the episode:


$\pi(s) \leftarrow \arg \max_a Q(s, a)$

**SLOW!**  
Assumption of infinite visits of  
all states and all actions!






# Reinforcement learning




# Reinforcement learning with Markov Decision Processes (MDP)



$R_{t+1}$  and  $S_{t+1}$  only depend on  $R_t$ ,  $S_t$  and  $A_t$

# Reinforcement learning with Markov Decision Processes (MDP)


( $S, A, R, T, \Pi$ )



$R_{t+1}$  and  $S_{t+1}$  only depend on  $R_t$ ,  $S_t$  and  $A_t$

- $S \equiv$  set of **states**: how the environment and the agent are
- $A \equiv$  set of **actions**: what the agent is allowed do  
 $A:S \mapsto S$
- $R \equiv$  **reward**: feedback on agent's action (immediate)  
 $R:S \mapsto \mathcal{R}$
- $\Pi \equiv$  **Policy**: agent's strategy: it defines the action to take at each state  
 $\Pi:S \mapsto A$

# Reinforcement learning



- $S \equiv$  set of **states**: how the environment and the agent are
- $A \equiv$  set of **actions**: what the agent is allowed to do
- $R \equiv$  **reward**: feedback on agent's action (immediate)
- $\Pi \equiv$  **Policy**: agent's strategy: it defines the action to take at each state

# RL with MDP: on-policy temporal difference

# RL with MDP: on-policy temporal difference: SARSA

( $\text{State}_t, \text{Action}_t, \text{Reward}_t, \text{State}_{t+1}, \text{Action}_{t+1}$ )

Sarsa (on-policy TD control) for estimating  $Q \approx q_*$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize  $S$

Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

Repeat (for each step of episode):

Take action  $A$ , observe  $R, S'$

Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until  $S$  is terminal

# RL with MDP: on-policy temporal difference method: SARSA

( $\text{State}_t, \text{Action}_t, \text{Reward}_t, \text{State}_{t+1}, \text{Action}_{t+1}$ )

Sarsa (on-policy TD control) for estimating  $Q \approx q_*$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize  $S$

Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

Repeat (for each step of episode):

Take action  $A$ , observe  $R, S'$

Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$$

$S \leftarrow S'; A \leftarrow A';$

until  $S$  is terminal

“stochastic gradient descent” on  $Q$

# Temporal Difference vs. Monte Carlo

- TD can learn before termination of episodes.
- TD can be used for either **non-episodic or episodic tasks**.
- The update depends on **single stochastic** transition  $\Rightarrow$  **lower variance**.
- Updates use bootstrapping  $\Rightarrow$  estimate has some **bias**.
- TD updates exploit the **Markov** property.
- MC learning must wait until the end of episodes.
- MC only works for **episodic tasks**.
- The update depends on a sequence of **many** stochastic transitions  $\Rightarrow$  much **larger variance**.
- **Unbiased** estimate.
- MC updates does not exploit the Markov property, hence it can be effective in **non-Markovian** environments.



# Contents

1. Major types of eye movements
2. What triggers saccades?
3. Reinforcement learning
  1. Markov decision processes
  2. Q learning with Monte Carlo and SARSA
4. Eye movements to learn to solve visually guided tasks

# Possible add-ons

- $\epsilon$ -greedy policy


$$\Pi^*(s) = \begin{cases} \operatorname{argmax}_a Q(s,a) & \text{with probability } \epsilon \\ \text{random } a & \text{with probability } 1-\epsilon \end{cases}$$

- States only partially observable (thus there is uncertainty on the state representation)
- Highly dimensional states
- Multitask learning

# Break

- Questions?


# Problem formulation (S, A, R, $\Pi$ )



Tasks/behaviors:

1. avoid obstacle
2. sidewalk following
3. pick litter

# Problem formulation (S, A, R, $\Pi$ )




Tasks/behaviors:

1. avoid obstacle
2. sidewalk following
3. pick litter

- $S \equiv$  set of states (position in the environment angle vs distance). The state is estimated in function of the visual perception (the eyes of the agent)
- $A \equiv$  set of actions (turn -15deg, no turn, turn 15 deg. Walking and action selection at fixed rate)
- $R(s) \equiv$  reward (-4 if hit obstacle, +1 in sidewalk, +2 if pick litter)
- $\Pi(s) \equiv$  Policy ( $\epsilon$ -greedy)

# Problem formulation (S, A, R, $\Pi$ )



Each task has its 2D state representation:

1. distance and angle, relative to the agent, to the nearest obstacle:
2. angle of the center-line of the sidewalk relative to the agent and signed distance to the center of the sidewalk
3. distance and angle, relative to the agent, to the nearest litter

Tasks/behaviors:

1. avoid obstacle
2. sidewalk following
3. pick litter

# Modular RL for multi-tasking: behaviors

- Each behavior (sensory-action control) has the ability to direct the eye, perform appropriate visual processing to retrieve the information necessary for performance of the behavior's task, and choose an appropriate course of action.
- Each behavior is only allowed to access perception (thus move the eye) during 300ms. That behavior updates its state space using a Kalman filter, while the others propagate their estimates and track the uncertainties

# Modular RL for multi-tasking: behaviors

- Each behavior (sensory-action control) has the ability to direct the eye, perform appropriate visual processing to retrieve the information necessary for performance of the behavior's task, and choose an appropriate course of action.
- Each behavior is only allowed to access perception (thus move the eye) during 300ms. That behavior updates its state space using a Kalman filter, while the others propagate their estimates and track the uncertainties
- How to mediate between behaviors?



# Modular RL for multi-tasking with SARSA

**Single task**

**Multiple n-tasks**

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha(r + \gamma Q(s', a'))$$

$\Sigma$

# Modular RL for multi-tasking with SARSA

## Single task

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha(r + \gamma Q(s', a'))$$

## Multiple n-tasks

$$Q(s, a) \approx \sum_{i=1}^n Q_i(s_i, a)$$

$\Sigma$

# Modular RL for multi-tasking with SARSA

Q learning

## Single task

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha(r + \gamma Q(s', a'))$$

## Multiple n-tasks

$$Q(s, a) \approx \sum_{i=1}^n Q_i(s_i, a)$$

action selection

$$\Pi^*(s) = \begin{cases} \operatorname{argmax}_a Q(s, a) & \text{with probability } \epsilon \\ \text{random } a & \text{with probability } 1 - \epsilon \end{cases}$$

$\Sigma$

# Modular RL for multi-tasking with SARSA

Q learning

## Single task

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha(r + \gamma Q(s', a'))$$

## Multiple n-tasks

$$Q(s, a) \approx \sum_{i=1}^n Q_i(s_i, a)$$

action  
selection

$$\Pi^*(s) = \begin{cases} \operatorname{argmax}_a Q(s, a) & \text{with probability } \epsilon \\ \text{random } a & \text{with probability } 1 - \epsilon \end{cases}$$

$$\Pi^*(s) = \begin{cases} \operatorname{argmax}_a \sum Q_i(s, a) & \text{with probability } \epsilon \\ \text{random } a & \text{with probability } 1 - \epsilon \end{cases}$$




# Q-values and Policies

Obstacle Avoidance

Sidewalk following

Litter collection

Q-function






# Q-values and Policies

Obstacle Avoidance




Sidewalk following

Litter collection

Q-function



Policy



# How to select where to look?

Or how to select where to extract features to improve my state representation?

## Single behavior


1. For each possible eye movement estimate the **cost of the uncertainty** in the state representation if the movement is not made.
2. Select the **eye movement with highest** potential cost.

## Multiple n-behaviors


1. For each possible behavior, for each possible eye movement estimate the **cost of the uncertainty** in the state representation if the movement is not made
2. For each possible behavior, estimate the total cost
3. Select the behavior with larger total cost
4. Select the **eye movement with highest** potential cost.

# Eye movements

a)



b)




a) An overhead view of the virtual agent during seven time steps of the sidewalk navigation task.

b) State estimates during the same seven time steps. The light gray regions correspond to the 90% confidence bounds before any perception has taken place. When present, the black regions correspond to the 90% confidence bounds after an eye movement has been made.



# In summary

- 3 tasks
- Shared action space
- Non-shared state representation
- 2 controls: **body** and **eye**
- **body**: actions are selected to maximize reward by RL agent
- **eye**: Each behavior decides where to look/observe its state. Only one behavior is allowed at each time
- **Reward is maximized by moving eyes according to the behavior that loses the most if not observed.**






Tasks/behaviors:

1. avoid obstacle
2. sidewalk following
3. pick litter

# Summary

- Major types of eye movements
- What makes eyes move?
- Multi-task RL modeling visually guided behaviors



# Bibliography

- Krauzlis, Recasting the Smooth Pursuit Eye Movement System, 2004
- Bruce and Tsotsos, Saliency, attention, and visual search: An information theoretic approach, 2009
- Land, Eye movements and the control of actions in everyday life, 2006
- Hayhoe and Rothkop, Vision in the natural world, 2010
- Shinoda et al, What controls attention in natural environments, 2001

# Bibliography (cont)

- Hofsten and Rosander, The Development of Gaze Control and Predictive Tracking in Young Infants, 1995
- Li et al, Memory and visual search in naturalistic 2D and 3D environments, 2016
- Sprague and Ballard, Eye Movements for Reward Maximization, 2003
- Sutton and Barto, Reinforcement Learning: An Introduction, 2017, <http://incompleteideas.net/book/bookdraft2017nov5.pdf>

Extra slides